

Blaize Lahman

Madison, WI 53703 • blaizelahman@gmail.com • 954-870-9033 • <https://github.com/blaizerlahman> • <https://blaize.me>

Education

University of Wisconsin-Madison

Madison, WI

BS in Computer Science, BS in Data Science, and Certificate in Economic Analytics Student

May 2027

GPA: 3.68/4.0, Dean's List Fall 2024, Spring 2025, Fall 2025

Relevant Coursework: Artificial Intelligence, Agentic Development, Operating Systems, Data Structures and Algorithms, Object-Oriented Programming, Machine Learning and Statistical Modeling, Discrete Mathematics, Linear Algebra

Technical Skills and Projects

Proficient Languages: Python, Go, Typescript, Java, C, SQL, R

Tools: Claude Code, Git, AWS, PostgreSQL, Docker, Artifactory, Linux, React, Next.js, Spring Boot, pandas, NumPy, Scikit-learn

EnrollAlert - Course Enrollment Notification Service

May 2025 – July 2025

- Designed a full-stack cloud-native pipeline that tracks seat availability for over 5,000 UW-Madison courses and serves real-time data to 400+ active users via a REST API at enrollalert.com.
- Built a batched, concurrent scraper in Go to conduct scheduled scrapes of new data from the UW-Madison course enrollment API, ingesting 10k+ hits an hour and updating the PostgreSQL database with course seat changes.
- Deployed a TypeScript/Next.js powered frontend on Vercel that communicates with database via REST API endpoints to provide instant course search, course seat data retrieval, user authentication, and course alert management.
- Architected an AWS-based infrastructure to run automated scraping with Lambda/EventBridge and dispatch course alert notifications through SES, providing users with course data and alert emails within 10 minutes of seat changes.

College Football Prediction Model

May 2024 – August 2024

- Engineered a Python prediction system that trains a per-team ensemble model for all 138 FBS teams, ingesting play-by-play, betting lines, and SP+ data into a SQLite store to project point spreads for any matchup.
- Built a backtesting and grading engine that evaluates predictions against-the-spread across full seasons, bins results into confidence tiers by historical success rate, achieving 67% winner prediction accuracy across 875 games.
- Implemented spread prediction as a per-team regression on final scoring margin, tuning each Lasso model's L1 regularization strength via CV grid search to drive automatic feature selection, then wrapping the tuned estimator in a bagging ensemble to decrease variance.

Work Experience

Dell Technologies - Client Solutions Group

Austin, TX

Software Engineer Intern - Applied AI

May 2026 – August 2026

- Built a native containerized evaluation engine into a custom AI agentic harness. Ensured total data visibility to allow for data-driven insights on harness effectiveness and benchmarking against competitors on industry standard evaluations.
- Scaled existing evaluation infrastructure from one dedicated worker to distributed testing across entire worker fleet, bringing test throughput from 1-2 evals/hour to 30+/hour. Stood up and managed task repository in Artifactory to allow for prebuilt image fetching on distributed workers, decreasing average task startup time from 5-10 min to <1 min.
- Developed an internal agentic development benchmark, distilling 30+ enterprise-specific workflows into agentic tasks with Harbor and highlighting actionable areas of improvement. Implemented harness enhancements based on results to increase performance by 20%+ on internal and industry benchmarks.
- Designed and integrated a universal routing model to decrease harness costs by 60% while achieving success rates within ~5% of frontier models on internal/external benchmarks. Fully automated a classification training pipeline including conversation window capturing, LLM-as-a-judge labeling, and session-aware hysteresis.

UW-Madison Department of Computer Sciences

Madison, WI

Head Undergraduate Teaching Assistant - Data Programming II

January 2025 – Present

- Built and integrated a custom LLM feedback server into student Jupyter notebook homework problems, allowing for immediate, customized help based on current student work using a client/server model.
- Designed a lightweight notebook client that scans relevant student code for each question. Implemented by building a backward dependency slice with Python's symbol tables, following imports from student written modules, and capturing tracebacks/cell output, allowing for complete context and comprehensive LLM feedback.
- Deployed a containerized async API on a VM with Docker Compose, Caddy and MySQL. It handles 1000s of requests per project from 200+ students with per-student token auth and full request logging.
- Hosts regular office hours to personally assist students by clarifying course material and debugging Python code, catering to individual learning needs to strengthen understanding of data programming concepts and practices.
- Collaborates with head teaching assistant in leading lab sessions, guiding over 30 students through hands-on data programming exercises while creating an engaging and positive learning environment.